

“大数据工程”课程实验报告

题目：HDFS 的编程实验	学号姓名：21377061 范春	日期：2024.03.12
---------------	------------------	---------------

实验环境：

- 1、虚拟机软件：VMware
- 2、Hadoop 版本：3.1.3
- 3、Java 版本：Oracle JDK 1.8
- 4、Java IDE：Eclipse

实验内容与完成情况：

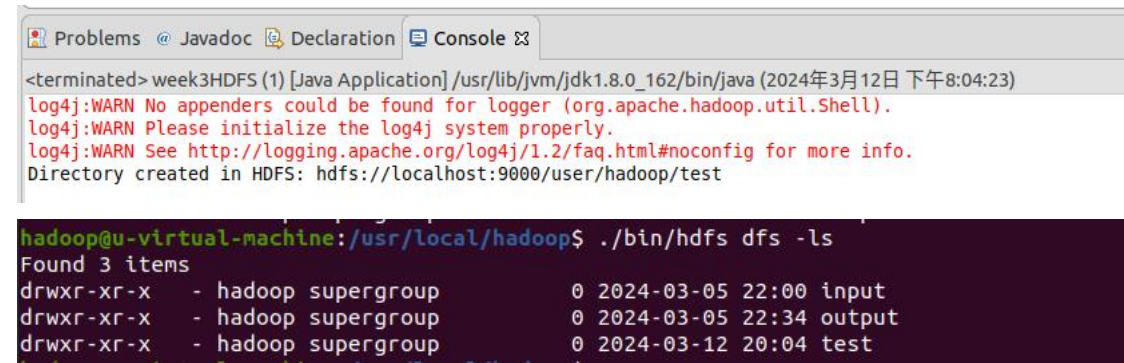
一、Eclipse 直接编译：

1、创建新的文件目录

```
public static void mkdir(String dirName) {
    try {
        Configuration conf = new Configuration();
        conf.set("fs.defaultFS", "hdfs://localhost:9000");
        conf.set("fs.hdfs.impl", "org.apache.hadoop.hdfs.DistributedFileSystem");
        FileSystem fs = FileSystem.get(conf);
        Path dirPath = new Path(dirName);

        if (fs.exists(dirPath)) {
            System.out.println("Directory already exists in HDFS.");
        } else {
            fs.mkdirs(dirPath);
            System.out.println("Directory created in HDFS: " + dirName);
        }

        fs.close();
    } catch (Exception e) {
        e.printStackTrace();
    }
}
```



2、创建新的文件

```
public static void createFile(String fileName) {
```

```

try {
    Configuration conf = new Configuration();
    conf.set("fs.defaultFS", "hdfs://localhost:9000");
    conf.set("fs.hdfs.impl", "org.apache.hadoop.hdfs.DistributedFileSystem");
    FileSystem fs = FileSystem.get(conf);
    Path filePath = new Path(fileName);

    if (!fs.exists(filePath.getParent())) {
        mkdir(filePath.getParent().toString());
    }

    fs.create(filePath).close();
    System.out.println("File created in HDFS: " + fileName);

    fs.close();
} catch (Exception e) {
    e.printStackTrace();
}
}

```

The screenshot shows an IDE interface with a 'Console' tab. The console output includes log4j warnings: 'WARN No appenders could be found for logger (org.apache.hadoop.util.Shell).', 'WARN Please initialize the log4j system properly.', and 'WARN See http://logging.apache.org/log4j/1.2/faq.html#noconfig for more info.'. Below these, it states 'File created in HDFS: hdfs://localhost:9000/user/hadoop/test/test.txt'. The terminal window below shows the command 'hadoop@u-virtual-machine:/usr/local/hadoop\$./bin/hdfs dfs -ls /user/hadoop/test' and its output: 'Found 1 items' followed by a table of file details for 'test.txt'.

```

<terminated> week3HDFS (1) [Java Application] /usr/lib/jvm/jdk1.8.0_162/bin/java (2024年3月12日 下午8:05:30)
log4j:WARN No appenders could be found for logger (org.apache.hadoop.util.Shell).
log4j:WARN Please initialize the log4j system properly.
log4j:WARN See http://logging.apache.org/log4j/1.2/faq.html#noconfig for more info.
File created in HDFS: hdfs://localhost:9000/user/hadoop/test/test.txt

hadoop@u-virtual-machine:/usr/local/hadoop$ ./bin/hdfs dfs -ls /user/hadoop/test
Found 1 items
-rw-r--r--  3 hadoop supergroup          0 2024-03-12 20:05 /user/hadoop/test/test.txt

```

3、删除文件

```

public static void delFile(String fileName) {
    try {
        Configuration conf = new Configuration();
        conf.set("fs.defaultFS", "hdfs://localhost:9000");
        conf.set("fs.hdfs.impl", "org.apache.hadoop.hdfs.DistributedFileSystem");
        FileSystem fs = FileSystem.get(conf);
        Path filePath = new Path(fileName);

        if (!fs.exists(filePath)) {
            throw new Exception("File does not exist in HDFS: " + fileName);
        }

        fs.delete(filePath, true);
        System.out.println("File deleted in HDFS: " + fileName);

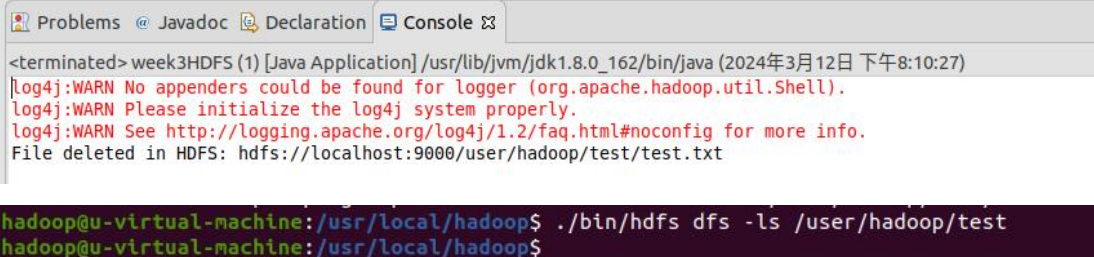
        fs.close();
    }
}

```

```

    } catch (Exception e) {
        e.printStackTrace();
    }
}

```



The screenshot shows an IDE interface. The top part is a console window with the following output:

```

<terminated> week3HDFS (1) [Java Application] /usr/lib/jvm/jdk1.8.0_162/bin/java (2024年3月12日 下午8:10:27)
log4j:WARN No appenders could be found for logger (org.apache.hadoop.util.Shell).
log4j:WARN Please initialize the log4j system properly.
log4j:WARN See http://logging.apache.org/log4j/1.2/faq.html#noconfig for more info.
File deleted in HDFS: hdfs://localhost:9000/user/hadoop/test/test.txt

```

The bottom part is a terminal window showing the following commands and output:

```

hadoop@u-virtual-machine: /usr/local/hadoop$ ./bin/hdfs dfs -ls /user/hadoop/test
hadoop@u-virtual-machine: /usr/local/hadoop$

```

4、自动批量创建文件并获取文件相关元信息

```

public static void batchCreateReadFiles(String dirName) {
    try {
        Configuration conf = new Configuration();
        conf.set("fs.defaultFS", "hdfs://localhost:9000");
        conf.set("fs.hdfs.impl", "org.apache.hadoop.hdfs.DistributedFileSystem");
        FileSystem fs = FileSystem.get(conf);
        Path dirPath = new Path(dirName);

        long startTime = Instant.now().getEpochSecond();
        long duration = 60; // 1 minute

        while (Instant.now().getEpochSecond() - startTime < duration) {
            String fileName = dirName + "/" + Instant.now().getEpochSecond() +
".log";

            fs.create(new Path(fileName)).close();
            System.out.println("File created in HDFS: " + fileName);

            Thread.sleep(10000); // 10 seconds
        }

        FileStatus[] fileStatuses = fs.listStatus(dirPath);
        for (FileStatus status : fileStatuses) {
            System.out.println("File Path: " + status.getPath());
            System.out.println("File Size: " + status.getLen() + " bytes");
            System.out.println("File Permissions: " + status.getPermission());
            System.out.println("File Creation Time: " + status.getModificationTime());
            System.out.println();
        }

        fs.close();
    } catch (Exception e) {
        e.printStackTrace();
    }
}

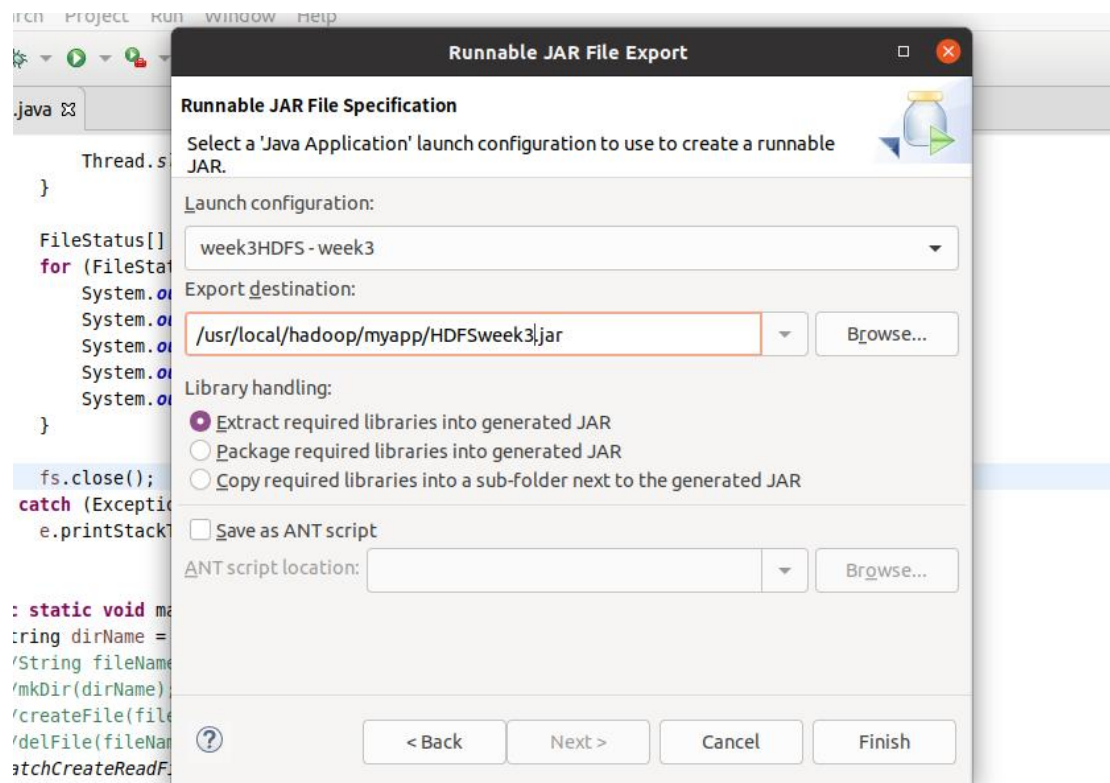
```

```
}  
}
```

```
hadoop@u-virtual-machine:/usr/local/hadoop$ ./bin/hdfs dfs -ls /user/hadoop/test  
Found 6 items  
-rw-r--r-- 3 hadoop supergroup 0 2024-03-12 20:22 /user/hadoop/test/1710246129.log  
-rw-r--r-- 3 hadoop supergroup 0 2024-03-12 20:22 /user/hadoop/test/1710246139.log  
-rw-r--r-- 3 hadoop supergroup 0 2024-03-12 20:22 /user/hadoop/test/1710246149.log  
-rw-r--r-- 3 hadoop supergroup 0 2024-03-12 20:22 /user/hadoop/test/1710246159.log  
-rw-r--r-- 3 hadoop supergroup 0 2024-03-12 20:22 /user/hadoop/test/1710246169.log  
-rw-r--r-- 3 hadoop supergroup 0 2024-03-12 20:22 /user/hadoop/test/1710246179.log
```

```
Problems @ Javadoc Declaration Console  
<terminated> week3HDFS (1) [Java Application] /usr/lib/jvm/jdk1.8.0_162/bin/java (2024年3月12日 下午8:22:08)  
log4j:WARN No appenders could be found for logger (org.apache.hadoop.util.Shell).  
log4j:WARN Please initialize the log4j system properly.  
log4j:WARN See http://logging.apache.org/log4j/1.2/faq.html#noconfig for more info.  
File created in HDFS: hdfs://localhost:9000/user/hadoop/test/1710246129.log  
File created in HDFS: hdfs://localhost:9000/user/hadoop/test/1710246139.log  
File created in HDFS: hdfs://localhost:9000/user/hadoop/test/1710246149.log  
File created in HDFS: hdfs://localhost:9000/user/hadoop/test/1710246159.log  
File created in HDFS: hdfs://localhost:9000/user/hadoop/test/1710246169.log  
File created in HDFS: hdfs://localhost:9000/user/hadoop/test/1710246179.log  
File Path: hdfs://localhost:9000/user/hadoop/test/1710246129.log  
File Size: 0 bytes  
File Permissions: rw-r--r--  
File Creation Time: 1710246129224  
  
File Path: hdfs://localhost:9000/user/hadoop/test/1710246139.log  
File Size: 0 bytes  
File Permissions: rw-r--r--  
File Creation Time: 1710246139237  
  
File Path: hdfs://localhost:9000/user/hadoop/test/1710246149.log  
File Size: 0 bytes  
File Permissions: rw-r--r--  
File Creation Time: 1710246149244  
  
File Path: hdfs://localhost:9000/user/hadoop/test/1710246159.log  
File Size: 0 bytes  
File Permissions: rw-r--r--  
File Creation Time: 1710246159252  
  
File Path: hdfs://localhost:9000/user/hadoop/test/1710246169.log  
File Size: 0 bytes  
File Permissions: rw-r--r--  
File Creation Time: 1710246169261  
  
File Path: hdfs://localhost:9000/user/hadoop/test/1710246179.log  
File Size: 0 bytes  
File Permissions: rw-r--r--  
File Creation Time: 1710246179269
```


二、将 Java 应用程序生成 JAR 包，部署到 Hadoop 平台上进行运行



```
hadoop@u-virtual-machine:/usr/local/hadoop$ mkdir myapp
hadoop@u-virtual-machine:/usr/local/hadoop$ cd /usr/local/hadoop/myapp
hadoop@u-virtual-machine:/usr/local/hadoop/myapp$ ls
HDFSweek3.jar
```

```
hadoop@u-virtual-machine:/usr/local/hadoop$ ./bin/hadoop jar ./myapp/HDFSweek3.jar
File created in HDFS: hdfs://localhost:9000/user/hadoop/test/1710249314.log
File created in HDFS: hdfs://localhost:9000/user/hadoop/test/1710249324.log
File created in HDFS: hdfs://localhost:9000/user/hadoop/test/1710249334.log
File created in HDFS: hdfs://localhost:9000/user/hadoop/test/1710249344.log
File created in HDFS: hdfs://localhost:9000/user/hadoop/test/1710249354.log
File created in HDFS: hdfs://localhost:9000/user/hadoop/test/1710249364.log
File Path: hdfs://localhost:9000/user/hadoop/test/1710249314.log
File Size: 0 bytes
File Permissions: rw-r--r--
File Creation Time: 1710249314212

File Path: hdfs://localhost:9000/user/hadoop/test/1710249324.log
File Size: 0 bytes
File Permissions: rw-r--r--
File Creation Time: 1710249324222

File Path: hdfs://localhost:9000/user/hadoop/test/1710249334.log
File Size: 0 bytes
File Permissions: rw-r--r--
File Creation Time: 1710249334230

File Path: hdfs://localhost:9000/user/hadoop/test/1710249344.log
File Size: 0 bytes
File Permissions: rw-r--r--
File Creation Time: 1710249344237

File Path: hdfs://localhost:9000/user/hadoop/test/1710249354.log
File Size: 0 bytes
File Permissions: rw-r--r--
File Creation Time: 1710249354244

File Path: hdfs://localhost:9000/user/hadoop/test/1710249364.log
File Size: 0 bytes
File Permissions: rw-r--r--
File Creation Time: 1710249364254
hadoop@u-virtual-machine:/usr/local/hadoop$
```

出现的问题：导出 jar 包时，不知道为什么 launch configuration 选项下有两个选项（week3HDFS-week3 和 week3HDFS-week3（1）），我一开始选择了 week3HDFS-week3，导出的 jar 包在运行时出现了报错：RunJar jarFile [mainClass] args...以为是技术问题，寻找了很多方法依然没有解决，后来重新选择 week3HDFS-week3（1）导出 jar 包就能正常运行了。

完整源码如下所示：

```
package week3;

import org.apache.hadoop.conf.Configuration;
import org.apache.hadoop.fs.*;
import java.time.Instant;

public class week3HDFS {
    public static void mkDir(String dirName) {
        try {
            Configuration conf = new Configuration();
            conf.set("fs.defaultFS", "hdfs://localhost:9000");
            conf.set("fs.hdfs.impl", "org.apache.hadoop.hdfs.DistributedFileSystem");
            FileSystem fs = FileSystem.get(conf);
            Path dirPath = new Path(dirName);

            if (fs.exists(dirPath)) {
                System.out.println("Directory already exists in HDFS.");
            } else {
                fs.mkdirs(dirPath);
                System.out.println("Directory created in HDFS: " + dirName);
            }

            fs.close();
        } catch (Exception e) {
            e.printStackTrace();
        }
    }

    public static void createFile(String fileName) {
        try {
            Configuration conf = new Configuration();
            conf.set("fs.defaultFS", "hdfs://localhost:9000");
            conf.set("fs.hdfs.impl", "org.apache.hadoop.hdfs.DistributedFileSystem");
            FileSystem fs = FileSystem.get(conf);
            Path filePath = new Path(fileName);

            if (!fs.exists(filePath.getParent())) {
                mkDir(filePath.getParent().toString());
            }
        }
    }
}
```

```

    }

    fs.create(filePath).close();
    System.out.println("File created in HDFS: " + fileName);

    fs.close();
} catch (Exception e) {
    e.printStackTrace();
}
}

public static void delFile(String fileName) {
    try {
        Configuration conf = new Configuration();
        conf.set("fs.defaultFS", "hdfs://localhost:9000");
        conf.set("fs.hdfs.impl", "org.apache.hadoop.hdfs.DistributedFileSystem");
        FileSystem fs = FileSystem.get(conf);
        Path filePath = new Path(fileName);

        if (!fs.exists(filePath)) {
            throw new Exception("File does not exist in HDFS: " + fileName);
        }

        fs.delete(filePath, true);
        System.out.println("File deleted in HDFS: " + fileName);

        fs.close();
    } catch (Exception e) {
        e.printStackTrace();
    }
}

public static void batchCreateReadFiles(String dirName) {
    try {
        Configuration conf = new Configuration();
        conf.set("fs.defaultFS", "hdfs://localhost:9000");
        conf.set("fs.hdfs.impl", "org.apache.hadoop.hdfs.DistributedFileSystem");
        FileSystem fs = FileSystem.get(conf);
        Path dirPath = new Path(dirName);

        long startTime = Instant.now().getEpochSecond();
        long duration = 60; // 1 minute

        while (Instant.now().getEpochSecond() - startTime < duration) {

```

```

        String fileName = dirName + "/" + Instant.now().getEpochSecond() + ".log";
        fs.create(new Path(fileName)).close();
        System.out.println("File created in HDFS: " + fileName);

        Thread.sleep(10000); // 10 seconds
    }

    FileStatus[] fileStatuses = fs.listStatus(dirPath);
    for (FileStatus status : fileStatuses) {
        System.out.println("File Path: " + status.getPath());
        System.out.println("File Size: " + status.getLen() + " bytes");
        System.out.println("File Permissions: " + status.getPermission());
        System.out.println("File Creation Time: " + status.getModificationTime());
        System.out.println();
    }

    fs.close();
} catch (Exception e) {
    e.printStackTrace();
}
}

public static void main(String[] args) {
    String dirName = "hdfs://localhost:9000/user/hadoop/test";
    //String fileName = "hdfs://localhost:9000/user/hadoop/test/test.txt";
    //mkdir(dirName);
    //createFile(fileName);
    //delFile(fileName);
    batchCreateReadFiles(dirName);
}
}

```